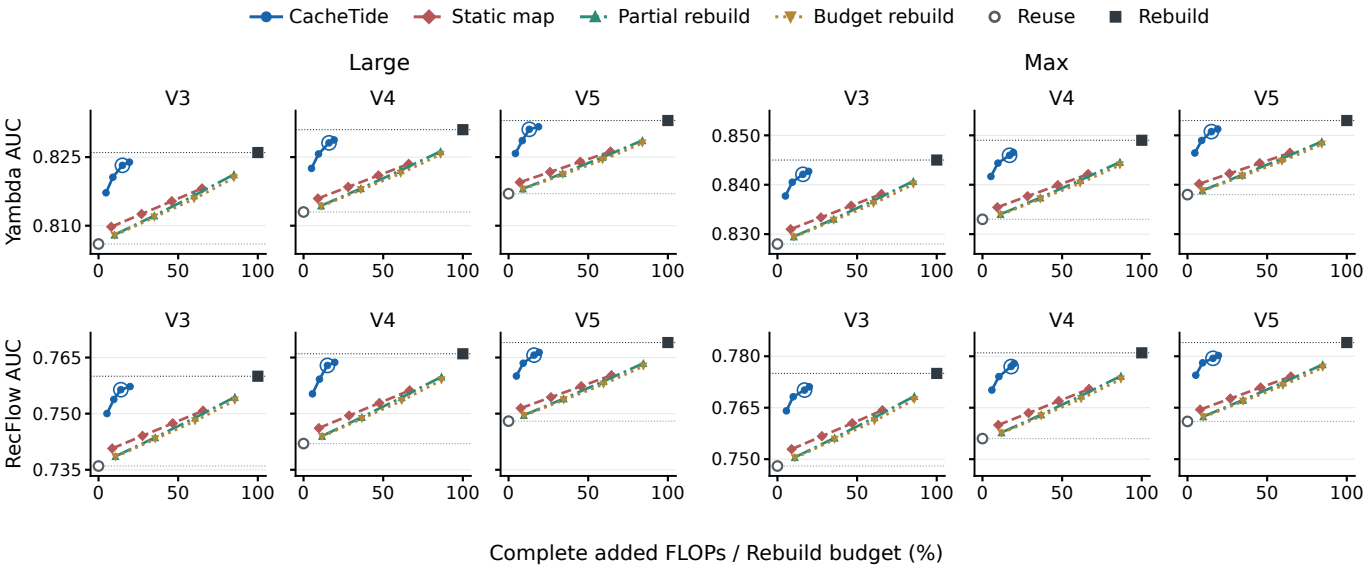
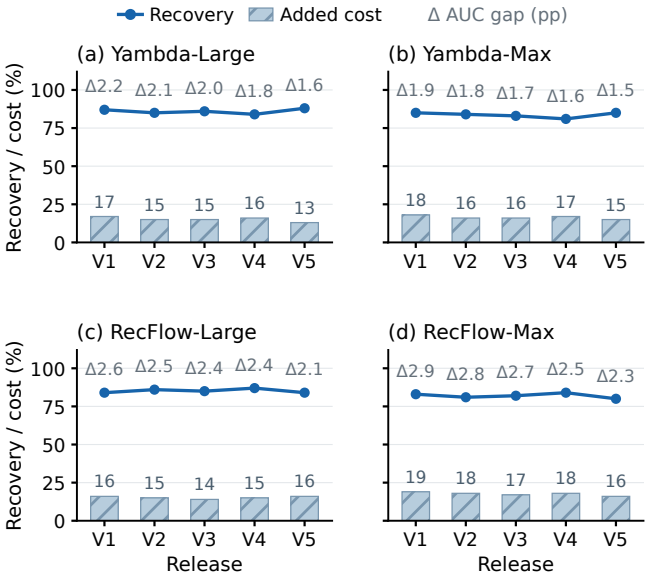
CacheTide | double-column main result and column-width details

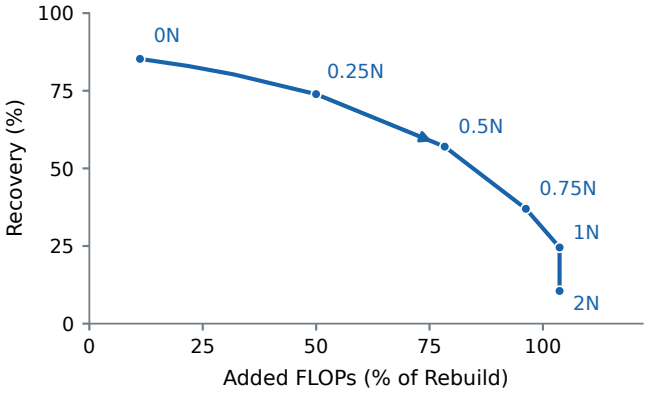
SIMULATED DATA - LAYOUT ONLY; NOT EXPERIMENTAL EVIDENCE



E1 (simulated). Twelve per-release quality versus complete-cost frontiers.



E3 (simulated). Default-point recovery with per-release quality gaps and costs.



E3b (simulated). One fixed release: quality and cumulative cost through two window turnovers.

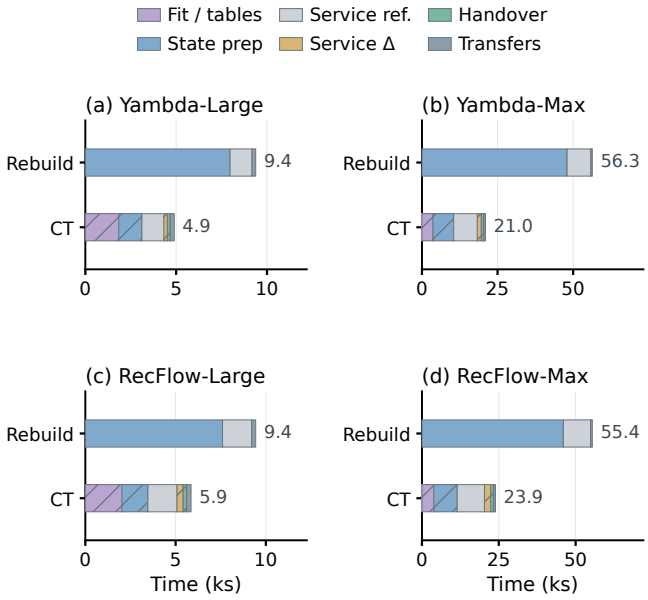
CacheTide | double-column main result and column-width details

SIMULATED DATA - LAYOUT ONLY; NOT EXPERIMENTAL EVIDENCE

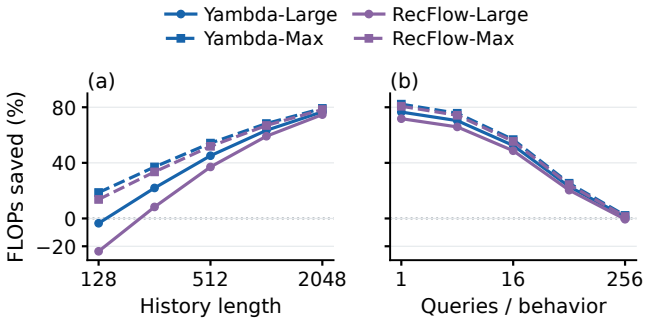
Configuration	Const. gain	Dynamic gain
Yambda-Large	1.40	1.55
Yambda-Max	1.23	1.33
RecFlow-Large	1.71	1.96
RecFlow-Max	1.85	2.05

Store/read	Gain (pp)	Read +%	Code %
36 / 6	1.19	12	5.62
72 / 18	1.62	32	11.25
72 / 6	1.55	12	11.25

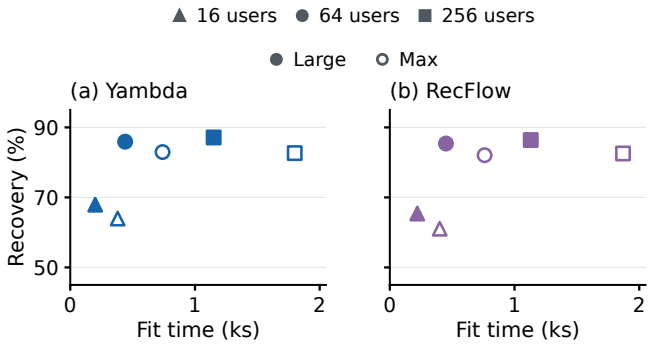
E4 (simulated). Table: candidate conditioning and code-width tradeoffs.



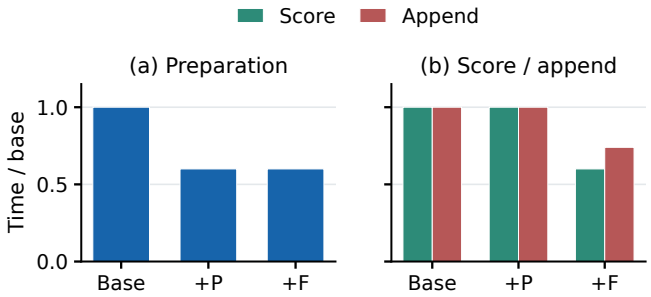
E2 (simulated). Complete processing-time attribution on matched resources.



E7 (simulated). Cost-model boundaries by history length and read density.



E5 (simulated). Discrete calibration quality--cost tradeoffs.



E6 (simulated). Defined cumulative projection and fusion optimizations.